

Affinify: AI-Powered Protein-Ligand Binding Affinity Predictor

An Interactive Platform for Drug Discovery Simulation

Project Synopsis

Lead Collaborator: Pranav Verma
Class: XII Aryabhatta
Subject: Computer Science / Biology
School: Lotus Valley International School
Project Duration: 8-10 Weeks
Date: July 7, 2025



Contents

1	Abstract	3
2	Introduction and Background	3
2.1	Problem Statement	3
2.2	Motivation	3
2.3	Scientific Relevance	3
2.4	Current State of Drug Discovery	4
3	Objectives	4
3.1	Primary Objectives	4
3.2	Secondary Objectives	4
4	Methodology	4
4.1	Project Architecture	4
4.2	Data Collection and Processing	5
4.2.1	Data Sources	5
4.2.2	Data Preprocessing	5
4.3	Feature Engineering	5
4.3.1	Molecular Descriptors	5
4.3.2	Protein Features	5
4.4	Machine Learning Models	5
4.4.1	Traditional Machine Learning	5
4.4.2	Deep Learning Approaches	5
5	Implementation Plan	6
5.1	Phase 1: Environment Setup and Data Collection (Weeks 1-2)	6
5.2	Phase 2: Data Processing and Feature Engineering (Weeks 3-4)	6
5.3	Phase 3: Model Development (Weeks 5-6)	6
5.4	Phase 4: Visualization and Web Interface (Weeks 7-8)	6
5.5	Phase 5: Testing and Expo Preparation (Weeks 9-10)	7
6	Expected Outcomes	7
6.1	Technical Deliverables	7
6.2	Educational Impact	7
6.3	Performance Metrics	8
7	Resources Required	8
7.1	Hardware Requirements	8
7.2	Software and Libraries	8
7.3	Data Resources	8
8	Risk Assessment and Mitigation	8
8.1	Technical Risks	8
8.2	Timeline Risks	9
9	Innovation and Uniqueness	9

10 Conclusion

9

1 Abstract

This project presents the development of an AI-powered platform for predicting protein-ligand binding affinity, a critical component in modern drug discovery. The system combines machine learning algorithms with molecular visualization techniques to create an interactive tool that can predict how strongly potential drug compounds will bind to target proteins. The platform features a user-friendly web interface built using Streamlit, 3D molecular visualization capabilities, and multiple predictive models including traditional machine learning and deep learning approaches.

The primary objective of this project is to demonstrate the application of artificial intelligence in pharmaceutical research while creating an educational tool that simplifies complex biochemical concepts for students and researchers. Specifically, the project aims to achieve:

- Prediction accuracy with $R^2 > 0.7$
- Processing speeds under 7 seconds per prediction

These goals ensure the platform is suitable for real-time demonstrations and educational purposes.

2 Introduction and Background

2.1 Problem Statement

Drug discovery is a complex, time-consuming, and expensive process that traditionally takes 10-15 years and costs billions of dollars. One of the most critical steps in this process is identifying how strongly potential drug molecules (ligands) bind to their target proteins. Current computational methods are either too simplistic or require extensive computational resources, making them inaccessible for educational purposes.

2.2 Motivation

The integration of artificial intelligence in drug discovery has revolutionized pharmaceutical research. By developing an interactive platform that demonstrates these concepts, we can:

- Showcase the practical applications of AI in healthcare
- Provide hands-on experience with molecular modeling and machine learning
- Create an educational tool for understanding protein-drug interactions
- Demonstrate the potential of computational biology in solving real-world problems

2.3 Scientific Relevance

Protein-ligand binding affinity prediction is fundamental to:

- Understanding disease mechanisms
- Designing more effective medications

- Reducing side effects through targeted therapy
- Accelerating the drug development pipeline
- Enabling personalized medicine approaches

2.4 Current State of Drug Discovery

Traditional drug discovery relies heavily on experimental screening of thousands of compounds, which is both time-intensive and costly. Recent advances in computational methods have introduced virtual screening techniques that can significantly reduce the number of compounds that need to be tested experimentally.

3 Objectives

3.1 Primary Objectives

1. Develop an AI-powered system for predicting protein-ligand binding affinity
2. Create an interactive web-based platform for real-time predictions
3. Implement 3D molecular visualization capabilities
4. Achieve prediction accuracy suitable for educational demonstrations
5. Design user-friendly interfaces for non-technical users

3.2 Secondary Objectives

1. Compare performance of different machine learning algorithms
2. Create educational content explaining AI applications in drug discovery
3. Develop batch processing capabilities for multiple predictions
4. Implement data visualization tools for result analysis
5. Create comprehensive documentation and user guides

4 Methodology

4.1 Project Architecture

The project follows a modular architecture with distinct components:

- **Data Collection Module:** Automated downloading and processing of molecular datasets
- **Feature Engineering Pipeline:** Extraction of molecular descriptors and structural features
- **Machine Learning Models:** Implementation of multiple predictive algorithms
- **Visualization Engine:** 3D molecular rendering and interactive plots
- **Web Interface:** User-friendly platform for predictions and analysis

4.2 Data Collection and Processing

4.2.1 Data Sources

4.2.2 Data Preprocessing

- Cleaning and standardization of binding affinity values
- Removal of duplicate entries and outliers
- Conversion of SMILES strings to molecular descriptors
- Extraction of protein features from PDB files
- Generation of interaction fingerprints

4.3 Feature Engineering

4.3.1 Molecular Descriptors

- Physicochemical properties (molecular weight, logP, polar surface area)
- Topological descriptors (connectivity indices, shape descriptors)
- Electronic properties (charge distribution, electronegativity)
- 3D geometric features (volume, surface area, conformational flexibility)

4.3.2 Protein Features

- Amino acid composition and sequence properties
- Secondary structure content
- Binding pocket characteristics
- Electrostatic potential maps

4.4 Machine Learning Models

4.4.1 Traditional Machine Learning

- **Random Forest:** Ensemble method handling non-linear relationships
- **XGBoost:** Gradient boosting for high-performance predictions
- **Support Vector Regression:** Effective for high-dimensional molecular data

4.4.2 Deep Learning Approaches

- **Deep Neural Networks:** Multi-layer networks for complex pattern recognition
- **3D Convolutional Neural Networks:** Processing volumetric molecular representations
- **Attention Mechanisms:** Focusing on important molecular interactions

5 Implementation Plan

5.1 Phase 1: Environment Setup and Data Collection (Weeks 1-2)

- Setup development environment with required libraries
- Create project structure and version control
- Download and organize molecular datasets
- Implement data validation and quality checks
- Setup local database for efficient data storage

5.2 Phase 2: Data Processing and Feature Engineering (Weeks 3-4)

- Develop molecular data processing pipelines
- Implement feature extraction algorithms
- Create standardized data formats
- Generate training and testing datasets
- Validate feature engineering approaches

5.3 Phase 3: Model Development (Weeks 5-6)

- Implement traditional machine learning models
- Develop deep learning architectures
- Optimize hyperparameters using cross-validation
- Compare model performances
- Select best-performing models for deployment

5.4 Phase 4: Visualization and Web Interface (Weeks 7-8)

- Create 3D molecular visualization components
- Develop interactive plotting functions
- Build Streamlit web application
- Implement real-time prediction capabilities
- Design user-friendly interfaces

5.5 Phase 5: Testing and Expo Preparation (Weeks 9-10)

- Comprehensive testing and debugging
- Performance optimization
- Creation of demonstration scenarios
- Preparation of presentation materials
- Documentation and user guides

Week	Task	Milestone
1-2	Environment Setup & Data Collection	Functional development environment
3-4	Data Processing & Feature Engineering	Processed datasets ready
5-6	Model Development	Trained ML models
7-8	Visualization & Web Interface	Functional web application
9-10	Testing & Expo Preparation	Complete demonstration ready

Table 1: Project Timeline and Key Milestones

6 Expected Outcomes

6.1 Technical Deliverables

- Functional AI models with $R^2 \geq 0.7$ prediction accuracy
- Interactive web application with real-time predictions
- 3D molecular visualization system
- Comprehensive documentation and tutorials
- Comparative analysis of different ML approaches

6.2 Educational Impact

- Demonstration of AI applications in healthcare
- Interactive learning tool for molecular biology concepts
- Hands-on experience with machine learning workflows
- Understanding of computational drug discovery processes
- Inspiration for careers in computational biology and AI

6.3 Performance Metrics

- **Accuracy:** $R^2 > 0.7$, RMSE < 1.0 pKd units
- **Speed:** < 5 seconds per prediction
- **Coverage:** 10,000+ protein-ligand pairs in the database
- **Usability:** Intuitive interface requiring no technical expertise

7 Resources Required

7.1 Hardware Requirements

- Computer with minimum 8GB RAM (16GB preferred)
- Graphics card for deep learning acceleration (required) **OR** Google Colab T4 GPU model (free access)
- Stable internet connection for data downloading
- Storage space for molecular datasets (minimum 50GB)

7.2 Software and Libraries

- Python 3.9+ with Anaconda distribution
- Machine learning libraries (TensorFlow, PyTorch, scikit-learn)
- Molecular processing tools (RDKit, BioPython)
- Visualization libraries (py3Dmol, Plotly, Streamlit)
- Development tools (Jupyter, Git, VS Code)

7.3 Data Resources

- BindingDB dataset (free access)
- PDBind dataset (registration required)
- ChEMBL database (API access)
- Protein Data Bank structures (free access)

8 Risk Assessment and Mitigation

8.1 Technical Risks

- **Data Quality Issues:** Mitigation through multiple data sources and validation
- **Model Performance:** Backup with simpler, proven algorithms
- **Computational Complexity:** Cloud computing resources if needed
- **Integration Challenges:** Modular development approach

8.2 Timeline Risks

- **Scope Creep:** Focus on core functionality first
- **Technical Difficulties:** Buffer time built into schedule
- **Data Access Delays:** Parallel development approach

9 Innovation and Uniqueness

This project offers several innovative aspects that distinguish it from existing solutions: the integration of multiple machine learning approaches within a single unified platform, real-time 3D visualization capabilities for protein-ligand interactions, an educational focus with simplified user interfaces accessible to non-technical users, a comprehensive comparative analysis framework for evaluating different algorithms, and interactive demonstration capabilities specifically designed for science expo presentations. The educational value extends beyond technical implementation by bridging computer science and biology curricula, demonstrating practical AI applications in healthcare, providing hands-on experience with cutting-edge technologies, encouraging student interest in computational biology careers, and creating a reusable educational resource for future learning and research endeavors.

10 Conclusion

This project represents an ambitious yet achievable integration of artificial intelligence, molecular biology, and web development technologies. By creating an interactive platform for protein-ligand binding affinity prediction, we aim to demonstrate the transformative potential of AI in healthcare while providing valuable educational experiences.

The project's success will be measured not only by technical metrics but also by its ability to inspire understanding of computational biology and showcase the practical applications of machine learning in solving real-world problems. The resulting platform will serve as both a functional tool for molecular analysis and an educational resource for future students and researchers.

Through careful planning, modular development, and focus on core functionality, this project has strong potential to be completed successfully within the allocated timeline while making a meaningful contribution to science education and demonstration of AI capabilities in the biological sciences.

Signature: _____

Student Name: Pranav Verma

Date: July 7, 2025

Signature: _____

Student Name: Ekaksh Goyal

Date: July 7, 2025

Legal Notice and Licensing Terms

Copyright Notice:

Copyright 2025 Pranav Verma, Ekaksh Goyal

Permission is granted to view this source code for educational or informational purposes only.

No permission is granted to reproduce, distribute, modify, or use the code or its underlying ideas in any way, in whole or in part, without explicit written permission from the authors.

This code is NOT open source.

Any commercial use, distribution, or reproduction requires prior written consent from the authors.

Teacher's Approval

I have read this project synopsis and find it to be a good project. I approve **Pranav Verma** and **Ekaksh Goyal** to work on this project: "**AI-Powered Protein-Ligand Binding Affinity Predictor**".

Teacher's Name: Ashu Arora
School: Lotus Valley International School
Signature: _____
Date: July 7, 2025